

Student Led R Session 03

Felix Mueller

3/21/2017

Topics of today

- ▶ writing / reading files with specific format
- ▶ CF 3 recommendation engines
- ▶ ggplot2
- ▶ handling NAs
- ▶ Feature derivation
- ▶ kNN
- ▶ naive Bayes
- ▶ trees
- ▶ Random Forests
- ▶ Logistic regression

writing / reading files with specific format

```
print(help(write.csv))
```

```
write.table(x, file = "", append = FALSE, quote = TRUE, sep = "  
", eol = "n", na = "NA", dec = ".", row.names = TRUE, col.names =  
TRUE, qmethod = c("escape", "double"), fileEncoding = "")
```

CF 3 recommendation engines I

Formulas from class (Introduction to Recommendation Engines, slide 59):

user similarity
(Pearson correlation)

$$\text{simil}(x, y) = \frac{\sum_{i \in I_{xy}} (r_{x,i} - \bar{r}_x)(r_{y,i} - \bar{r}_y)}{\sqrt{\sum_{i \in I_{xy}} (r_{x,i} - \bar{r}_x)^2 \sum_{i \in I_{xy}} (r_{y,i} - \bar{r}_y)^2}}$$

rating prediction
(weighted sum)

$$r_{u,i} = \bar{r}_u + k \sum_{u' \in U} \text{simil}(u, u') (r_{u',i} - \bar{r}_{u'}) \quad , k = 1 / \sum_{u' \in U} |\text{simil}(u, u')|$$

CF 3 recommendation engines II

```
victoria <- critics %>% filter(User=="Victoria") %>%  
  select(-User) %>% as.integer() #from 3  
similarity <- critics %>% filter(User != "Victoria"  
  ) %>% select(-User) %>% t() %>% cor(.,victoria,use=  
  "pairwise.complete.obs",method="pearson") #from 4  
avg_rating = critics %>% filter(User!="Victoria"  
  ) %>% select(-User)  
avg_rating=apply(avg_rating,1,function(x)mean(x,na.rm=T))  
avg_victoria = critics %>% filter(User=="Victoria") %>%  
  select(-User) %>% t() %>% mean(.,na.rm = T)
```

CF 3 recommendation engines III

```
critics_no_victoria = critics %>% filter(User !=  
  "Victoria") %>% select(-User)  
pred_by_movie = vector()  
for (i in seq(1:ncol(critics_no_victoria))){  
  pred_by_movie[i]=sum(similarity * (  
    critics_no_victoria[i]-avg_rating),na.rm=T)  
}  
#let's normalise the values  
rating_calculated <- avg_victoria + (  
  (1/sum(abs(similarity))) * pred_by_movie)  
movie_names <- names(critics[, -1])  
rating_calculated = as.data.frame(cbind(  
  rating_calculated, movie_names))
```

CF 3 recommendation engines IV

```
#exclude movies she has already watched
critics_victoria = critics %>% filter(User == "Victoria"
) %>% select(-User) %>% t() %>% is.na()
rating_calculated %>% filter(critics_victoria) %>%
  arrange(.,desc(rating_calculated)) %>% top_n(5,
rating_calculated)
```

	rating_calculated	movie_names
## 1	3.7917013044215	The Matrix
## 2	3.50776533175371	Forrest Gump
## 3	3.33118834864677	The Sixth Sense
## 4	3.11491825315719	Shakespeare in Love
## 5	2.9124513228665	Blade Runner

Possible kinds of feature engineering

- ▶ Adding values
- ▶ Multiplying values
- ▶ Decompose Categorical Attributes
- ▶ Decompose a Date-Time
- ▶ Categories to numeric and the other way around
- ▶ PCA
- ▶ ML models predicting as a variable (i.e. Bagged trees' predictions)

More Information: <http://machinelearningmastery.com/discover-feature-engineering-how-to-engineer-features-and-how-to-get-good-at-it/>

Decompose a Date-Time

```
median(df$date_recorded)
```

```
## [1] "2012-10-10"
```

```
max(df$date_recorded)
```

```
## [1] "2013-12-03"
```

```
min(df$date_recorded)
```

```
## [1] "2002-10-14"
```

```
df$new_date = if_else(df$date_recorded>median  
  (df$date_recorded), "recent", "not recent")  
head(df$new_date)
```

```
## [1] "not recent" "not recent" "not recent" "recent"
```

```
## [6] "not recent"
```

Handling NAs

- ▶ most similar value
- ▶ not use the feature
- ▶ create a new category for it
- ▶ not recommended: exclude instances from sample
- ▶ not recommended: do nothing

Handling NAs: applied I

#check for NAs

```
names = df[,apply(df, 2, function(x)
  any(is.na(x)))] %>% names()
apply(df[,names], 2, function(x) sum(is.na(x)))
```

##	funder	installer	subvillage
##	3269	3285	333
##	scheme_management	scheme_name	permit
##	3476	25424	2768

#we handle scheme_name NAs as a new category

#because NAs reflect about 50% of the data

```
index = sapply(df$scheme_name,is.na)
levels(df$scheme_name) = c(levels(df$scheme_name),
                           "unknown")
df$scheme_name[index] = "unknown"
```

Handling NAs: applied II

```
get_mode <- function(x) {  
  unique_x <- unique(x)  
  unique_x[which.max(tabulate(match(x, unique_x)))]  
}  
  
#update names  
names = df[,apply(df, 2, function(x)  
  any(is.na(x)))] %>% names()  
names
```

```
## [1] "funder"           "installer"        "subvillage"  
## [4] "public_meeting"   "scheme_management" "permit"
```

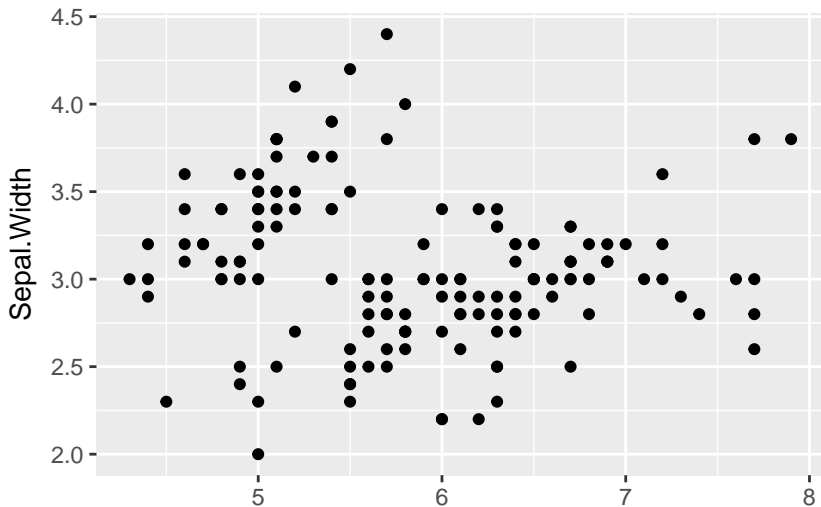
Handling NAs: applied III

```
df = data.frame(df)
for (name in names){
  mode = get_mode(df[,name])
  index = sapply(df[,name],is.na)
  levels(df[,name]) = c(levels(df[,name]),mode)
  df[index,name] = mode
}
#finished NA handling
print(df[,apply(df, 2, function(x) any(is.na(x)))]
      %>% names())
```

```
## character(0)
```

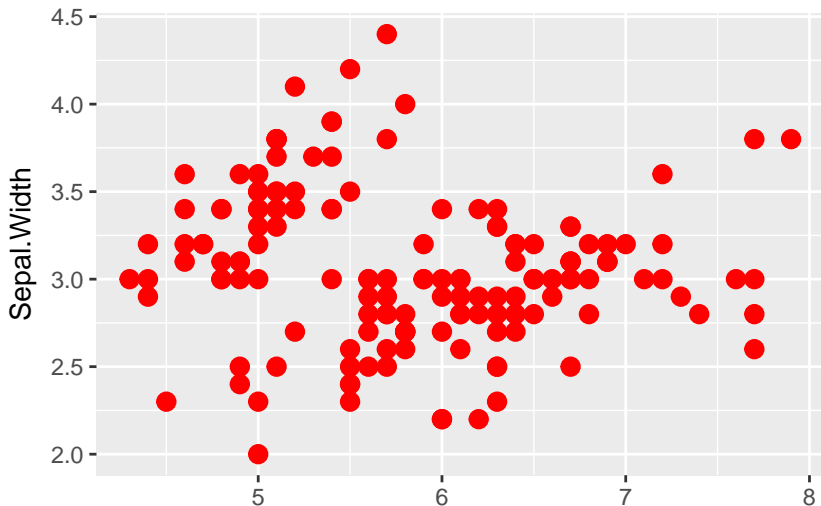
ggplot2 I

```
ggplot(data = iris, mapping =  
  aes(x=Sepal.Length,y=Sepal.Width))+  
  geom_point()
```



ggplot2 II

```
ggplot(data = iris, mapping =  
  aes(x=Sepal.Length,y=Sepal.Width))+  
  geom_point(colour = I('red'), size = 3)
```




ggplot2 III

```
ggplot(data = iris, mapping =  
  aes(x=Sepal.Length,y=Sepal.Width,  
  color=Species))+  
  geom_point()
```




ggplot2 IV

```
ggplot(data = iris, mapping =  
  aes(x=Species))+  
  geom_bar(colour=I("black"),fill=I("blue"),size=3)+  
  coord_flip()
```




ggplot2 V

```
d = diamonds %>% sample_n(100)
ggplot(data = d, mapping =
  aes(x=cut,y=price))+
  geom_boxplot(colour=I("green"))
```




ggplot2 VI

```
ggplot(data = d, mapping =  
  aes(carat, price, shape=cut)) +  
  geom_point()
```




ggplot2 VII

```
ggplot(data = d, mapping =  
  aes(carat,price,colour=cut))+  
  geom_point()
```




ggplot2 VIII

```
ggplot(data = d, mapping =  
  aes(carat,price,size=cut))+  
  geom_point()
```



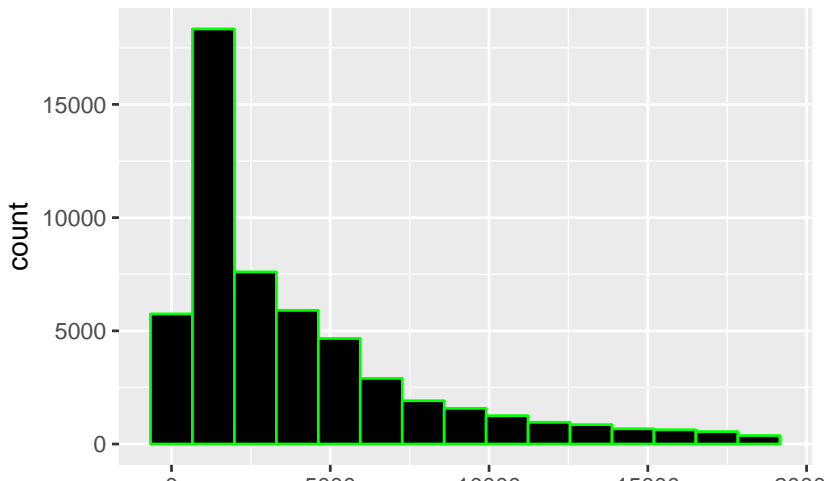
ggplot2 IX

```
ggplot(data = d, mapping =  
  aes(carat,price))+  
  geom_line(colour=I("blue"))+  
  geom_abline(intercept = median(d$price))
```



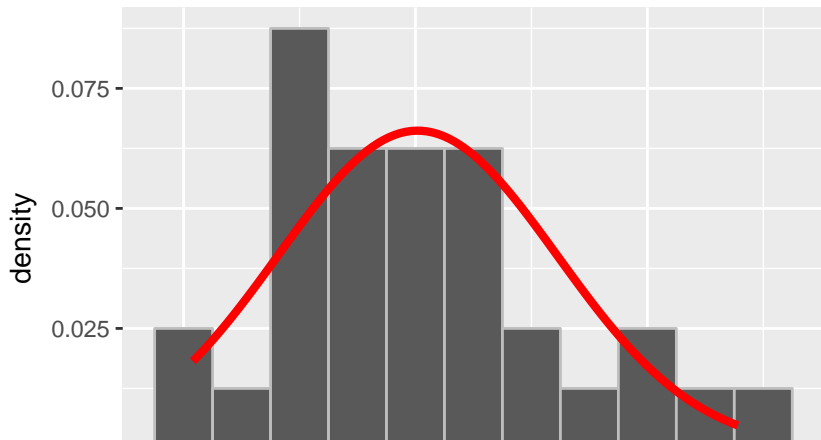
ggplot2 X

```
ggplot(data = diamonds, mapping =  
  aes(price)) +  
  geom_histogram(colour=I("green"), fill=I("black"),  
    bins = 15)
```



ggplot2 XI

```
ggplot(data=mtcars, aes(x=mpg)) +  
  geom_histogram(aes(y=..density..), binwidth=2.5,  
    colour=I("gray")) +  
  stat_function(fun=dnorm, args=list(mean=mean(mtcars$mpg),  
    sd=sd(mtcars$mpg)), colour=I("red"), size=1.5)
```



Your turn: model with the wine dataset

Input variables (based on physicochemical tests):

- ▶ fixed acidity
- ▶ volatile acidity
- ▶ citric acid
- ▶ residual sugar
- ▶ chlorides
- ▶ free sulfur dioxide
- ▶ total sulfur dioxide
- ▶ density
- ▶ pH
- ▶ sulphates
- ▶ alcohol

Output variable (based on sensory data):

- ▶ quality (score between 0 and 10)

How to download the wine dataset

```
wine <- read.csv2(url)
glimpse(wine)
```

```
## Observations: 4,898
```

```
## Variables: 12
```

```
## $ fixed.acidity      <fctr> 7, 6.3, 8.1, 7.2, 7.2, 8.1
```

```
## $ volatile.acidity  <fctr> 0.27, 0.3, 0.28, 0.23, 0.2
```

```
## $ citric.acid        <fctr> 0.36, 0.34, 0.4, 0.32, 0.3
```

```
## $ residual.sugar    <fctr> 20.7, 1.6, 6.9, 8.5, 8.5,
```

```
## $ chlorides          <fctr> 0.045, 0.049, 0.05, 0.058,
```

```
## $ free.sulfur.dioxide <fctr> 45, 14, 30, 47, 47, 30, 30
```

```
## $ total.sulfur.dioxide <fctr> 170, 132, 97, 186, 186, 97
```

```
## $ density           <fctr> 1.001, 0.994, 0.9951, 0.99
```

```
## $ pH                <fctr> 3, 3.3, 3.26, 3.19, 3.19,
```

```
## $ sulphates          <fctr> 0.45, 0.49, 0.44, 0.4, 0.4
```

```
## $ alcohol            <fctr> 8.8, 9.5, 10.1, 9.9, 9.9,
```

```
## $ quality            <int> 6, 6, 6, 6, 6, 6, 6, 6, 6,
```

kNN

Calculates the distance between data instances.

Good if:

- ▶ < 20 features
- ▶ many instances
- ▶ non-linear problems

Watch out if:

- ▶ too many features (Curse of multidimensionality)
- ▶ scale & center before training

naive Bayes

Computes the conditional probabilities of an event

Good if:

- ▶ missing data
- ▶ tons of features

Watch out if:

- ▶ highly correlated predictors
- ▶ use log probabilities for many features
- ▶ zero observations problem

trees

creates set of rules based on posterior & prior probabilities

Good if:

- ▶ categorical data
- ▶ interpretation necessary

Watch out if:

- ▶ many numeric attributes
- ▶ trees overfit your data → pruning!

Random Forest

creates several decorrelated decision trees (usually > 100)

Good if:

- ▶ categorical data
- ▶ correlated predictors
- ▶ big dataset

Watch out if:

- ▶ many numeric attributes

Logistic regression

estimates regression coefficients with Maximum Likelihood function

Good if:

- ▶ instances $k \gg$ predictors p
- ▶ used for feature selection with ~ 20 -50 predictors

Watch out if:

- ▶ too many predictive variables ($10k < p$)